

Neural Network and Deep Learning – ICP4

AVINASH BORRA

Student ID: 700756622

Github link: <https://github.com/Avinash-hub1/Assignment-4.git>

Video link: https://drive.google.com/file/d/1c3CaxBTwY-9SddgzuyuPV44KWqmVd8Dd/view?usp=drive_link

In class programming:

1. Follow the instruction below and then report how the performance changed.(apply all at once)

- Convolutional input layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
- Dropout layer at 20%.
- Convolutional layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
- Max Pool layer with size 2×2 .
- Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
- Dropout layer at 20%.
- Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
- Max Pool layer with size 2×2 .
- Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
- Dropout layer at 20%.
- Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
- Max Pool layer with size 2×2 .
- Flatten layer.
- Dropout layer at 20%.
- Fully connected layer with 1024 units and a rectifier activation function.
- Dropout layer at 20%.
- Fully connected layer with 512 units and a rectifier activation function.

- Dropout layer at 20%.
- Fully connected output layer with 10 units and a Softmax activation function

Did the performance change?

2. Predict the first 4 images of the test data using the above model. Then, compare with the actual label for those 4

images to check whether or not the model has predicted correctly.

3. Visualize Loss and Accuracy using the history object

Output:

```

import numpy as np
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D
from tensorflow.keras.constraints import MaxNorm
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.optimizers.schedules import ExponentialDecay
import matplotlib.pyplot as plt

# Fix random seed for reproducibility
seed = 7
np.random.seed(seed)

# Load data
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

# Normalize inputs from 0-255 to 0.0-1.0
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# One hot encode outputs
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
num_classes = y_test.shape[1]

# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))

```

2m 51s completed at 10:21 PM

Chrome File Edit View History Bookmarks Profiles Tab Window Help Thu Jul 11 22:38

colab.research.google.com/drive/1fGRCg-Cxs7Kc0GFYfaEVN3vhPUKu_TYN

Untitled3.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

# Compile model
epochs = 5
lr_rate = 0.01
lr_schedule = ExponentialDecay(
    initial_learning_rate=lr_rate,
    decay_steps=epochs * len(X_train) // 32,
    decay_rate=0.1
)
sgd = SGD(learning_rate=lr_schedule, momentum=0.9, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())

# Fit the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)

# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1] * 100))

# Predict the first 4 images of the test data
predictions = model.predict(X_test[:4])
predicted_classes = np.argmax(predictions, axis=1)
actual_classes = np.argmax(y_test[:4], axis=1)
```

2m 51s completed at 10:21 PM

Untitled3.ipynb - Colab New Tab

colab.research.google.com/drive/1fGRCg-Cxs7Kc0GFYfaEVN3vhPUKu_TYN

Untitled3.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
# Print the predictions and actual labels
print("Predicted classes: ", predicted_classes)
print("Actual classes: ", actual_classes)

# Plot the first 4 test images, predicted labels, and actual labels
fig, axes = plt.subplots(1, 4, figsize=(15, 3))
for i in range(4):
    axes[i].imshow(X_test[i])
    axes[i].set_title(f"Pred: {predicted_classes[i]}, Actual: {actual_classes[i]}")
    axes[i].axis('off')
plt.show()

# Visualize Loss and Accuracy
plt.figure(figsize=(12, 4))

# Plot Loss
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='train_loss')
plt.plot(history.history['val_loss'], label='val_loss')
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(loc='upper right')

# Plot Accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='train_accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(loc='upper left')

plt.show()
```

Download data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>

2m 51s completed at 10:21 PM

Chrome File Edit View History Bookmarks Profiles Tab Window Help Thu Jul 11 22:39

colab.research.google.com/drive/1fGRCg-Cxs7Kc0GfYfaEVN3vhPUKu_TYN

Untitled3.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

Download data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 [=====] - 4s 0us/step
Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
dropout (Dropout)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
dropout_1 (Dropout)	(None, 16, 16, 64)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
dropout_2 (Dropout)	(None, 8, 8, 128)	0
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dropout_3 (Dropout)	(None, 2048)	0
dense (Dense)	(None, 1024)	2098176

2m 51s completed at 10:21 PM

Untitled3.ipynb - Colab

colab.research.google.com/drive/1fGRCg-Cxs7Kc0GfYfaEVN3vhPUKu_TYN

Untitled3.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

dropout_4 (Dropout) (None, 1024) 0

dense_1 (Dense) (None, 512) 524800

dropout_5 (Dropout) (None, 512) 0

dense_2 (Dense) (None, 10) 5130

=====

Total params: 2915114 (11.12 MB)
Trainable params: 2915114 (11.12 MB)
Non-trainable params: 0 (0.00 Byte)

None

Epoch 1/5
1563/1563 [=====] - 28s 14ms/step - loss: 1.8439 - accuracy: 0.3196 - val_loss: 1.5040 - val_accuracy: 0.4444

Epoch 2/5
1563/1563 [=====] - 17s 11ms/step - loss: 1.4159 - accuracy: 0.4837 - val_loss: 1.3987 - val_accuracy: 0.4963

Epoch 3/5
1563/1563 [=====] - 16s 10ms/step - loss: 1.2221 - accuracy: 0.5603 - val_loss: 1.1673 - val_accuracy: 0.5711

Epoch 4/5
1563/1563 [=====] - 16s 10ms/step - loss: 1.0913 - accuracy: 0.6091 - val_loss: 1.0671 - val_accuracy: 0.6175

Epoch 5/5
1563/1563 [=====] - 14s 9ms/step - loss: 0.9978 - accuracy: 0.6435 - val_loss: 0.9806 - val_accuracy: 0.6526

Accuracy: 65.26%

1/1 [=====] - 0s 425ms/step

Predicted classes: [3 8 8 0]
Actual classes: [3 8 8 0]

Pred: 3, Actual: 3

Pred: 8, Actual: 8

Pred: 8, Actual: 8

Pred: 0, Actual: 0

2m 51s completed at 10:21 PM

